

Software Testing and Quality Assurance Final Project Report

Project: pandas - Python Data Analysis Toolkit

Team Members: Sandeep Ramavath, Nithikesh Reddy (Bobbili), Mallikarjuna

1. Project Overview

Pandas is an open-source library for data analysis in Python that offers quick and adaptable data structures for working with relational and labelled datasets. Pandas serves as a core component of any real-world analysis-related Python project.

1.1 Key Features

DataFrame and Series data structures for working with tabular data - Data cleaning and transformation capabilities for handling missing data - Powerful aggregation and grouping functionality for data analysis - I/O tools supporting CSV, Excel, SQL databases, HDF5, Parquet, and more - Time series functionality for working with date/time data

1.2 Project Scale

- Lines of Code: 484,492 lines of Python code
- Files: 7,218 Python files
- Usage: Used by millions of data scientists and analysts worldwide
- Domain: Data science, machine learning, financial analysis, scientific computing

Pandas is the chosen library for this extensive testing due to its: (1) Real-world Impact - an essential component of the Data Science Ecosystem (2) Complexity - a sufficiently large codebase to illustrate numerous testing approaches and techniques (3) Maturity - as a mature project, Pandas has an existing set of tests already created that can be used for a baseline of future enhancements. (4) Range of Functionality-Pandas operates on many types of data including Data I/O operations, manipulation and aggregation, and time series data.

2. Testing Stage Highlights

2.1 Unit Testing - Coverage Expansion

Key Takeaway: Edge case testing validates existing quality even when no bugs are found

We added fifteen more unit tests across three central modules in the pandas library, including: the nanops module with five tests, Series constructors with five tests, and five tests for DateTime offsets. In particular, the following tested: empty arrays with nan_sum, arrays with only NaN and nanargmax/nanargmin, boundary values for ddof, and error conditions in nanops; key errors in dictionaries, empty data, mixed types of data, and memory-intensive operations of

100,000 elements or greater in Series constructors; timezone edge conditions, daylight saving time transitions, and offset arithmetic in DateTime offsets.

The testing effort had a 100% pass rate, with all 15 tests passing. Coverage increased from 10% to 11%, an absolute gain of 1 percentage point. No defects were observed, confirming that pandas appropriately handles the examined edge cases.

The tests provide insight into the following: a) edge-case behavior related to empty and all-NaN arrays, b) memory management-patterns related to large Series operations, and c) timezone-related concerns and DST transitions of DateTime operations. Although no defects were uncovered, the tests provide useful regression protection and documentation of the expected behavior for these previously unconsidered edge cases.

2.2 Unit Testing - Mocking & Stubbing

A total of nine mocked tests were created, using 'pytest-mock' and 'monkeypatch' in order to achieve full test isolation, which includes the following: 1. **Database I/O tests** (3): 'read_sql' with connection failures; query errors; results empty (no records); 2. **File System I/O tests** (3): 'read_csv' or 'read_excel' when the files are not found; file contains corrupted data; file permissions not allowed; 3. **DateTime Operations** tests (3): 'Timestamp.now' and the 'date_range', using a fixed timestamp to determine if the tests will always return the same responses for a given test condition.

The tests confirmed full test isolation (all tests passed) and as such removed the impacts of time-series non-determinism; verified the behaviour of error handling during test conditions, which cannot be recreated with actual dependencies.

Mocking strategies were created for simulating database connection failure; techniques for simulating corrupted file data; and methods for generating reproducible time-dependent calculations. Thus, the tests created within this project can be employed to determine the error-handling capabilities of an application, without having to call on external resources (database server); corrupted files; or specific system time. Therefore the test battery will be more robust and better maintainable.

2.3 Mutation Testing

Key Conclusion: High test coverage does not inherently ensure that tests detect defects.

The following three validation functions within pandas/util/_validators.py were mutation tested:

- **validate_endpoints**: by Sandeep, 30 lines, validates the "closed" parameter for interval boundaries
- **validate_percentile** by Nithikesh: 30 lines, validates percentile values within the range [0, 1]
- **validate_bool_kwarg** (Mallikarjuna): validates boolean keyword arguments (43 lines)

Initial baseline: 20 random mutations per function

Improved Testing: 40 random mutations after improvements

Mutation Testing Detection Rates

Sandeep: validate_endpoints - Initial: 4/22 (18%); Final: 16/51 (31%)

Nithikesh: validate_percentile - Initial: 3/22 (14%); Final: 100% (in range)

Mallikarjuna: validate_bool_kwarg — Initial: 0/20 (0%); Final: 4/42 (10%)

Targeted tests were added to eliminate the surviving mutants, including:

- Enhancing test assertions to check the intended behavior, rather than test execution
- Expanding the boundary-condition testing Improving coverage of error paths This exercise shows that coverage metrics do not serve as a reliable indicator of test quality. Strong assertions and substantial edge-case testing are required for finding bugs effectively.

2.4 Static Analysis & Code Smells

Takeaway: Automated tools catch maintainability items prior to developing into bugs.

Dynamically scanning via ruff v0.14.3, Full Static Analysis on [pandas/util/_validators.py], All Checks Enabled.

There were 88 Code Smells discovered across multiple code smell categories. The following are all of the instances of each:

25 occurrences of F-String Literals in the Exception Messages (Similar to EM102)

23 occurrences of Missing Type Annotations (Similar to ANN001)

15 occurrences of Long Exception Messages (Similar to TRY003)

8 occurrences of Boolean Positional Arguments (Similar to FBT001 and FBT002)

1.EM102 Fix (Sandeep): Extracted f-string exception to variable

```
# Before
raise ValueError(f"{arg_name} must be...")

# After
msg = f"{arg_name} must be..."
raise ValueError(msg)
```

2.PLC0415 Fix (Nithikesh): Moved import to module top-level

```
# Before: Import inside function
def validate_fillna_kwargs(...):
    from pandas.core.missing import clean_fill_method

# After: Import at top
from pandas.core.missing import clean_fill_method
```

3.SIM102 Fix (Mallikarjuna): Combined nested if statements to reduce complexity

These fixes improve code maintainability, make testing easier, and follow Python best practices. The remaining 85 code smells are documented for future improvement.

2.5 Integration Testing

Key Finding: The interfaces between the modules of pandas are well-thought-out, with precise preservation of data types.

Six integration tests exercise interactions between different pandas modules:

Series-DataFrame-Dtype tests (2 tests): Test type preservation in `Series.to_frame()` and in DataFrame construction.

Validate-Missing Data - 2 tests: validate that `validate_fillna_kwargs` delegation correctly routes to the `clean_fill_method`.

Dtype Backend-Libs tests (2): test `check_dtype_backend` with special sentinel values `lib.no_default`.

All tests passed in 0.94 seconds with no defects found, confirming that:

- Strength of inter-module integration

Correct functioning of all type-preservation logic

- Proper Delegation within the Validation Framework

Specific findings included:

Series-to-DataFrame conversion maintains proper dtypes

DataFrame construction correctly handles mixed-dtype Series - Validation module correctly delegates to the missing data module - Backend validators handle C library sentinel values properly The lack of bugs shows that the module interfaces of pandas are properly designed, well-tested, and inter-component integration is sound and reliable.

2.6 System Testing

Key Takeaway:End-to-End Workflows and Data Integrity

This paper presents three different end-to-end workflows implemented using the publicly exposed pandas API; each workflow represents a complete user scenario, from data ingestion to export, and is validated for data integrity and correctness.

1) CSV I/O roundtrip (Sandeep)

Process: Loading a CSV file, data manipulation, and export back to CSV.

- Validation: The roundtrip preserves all data types and values.

2) Cleaning Missing Data (Nithikesh)

Process: Identify missing values, use relevant imputation techniques, and review the outcome.
Validation: Handling of missing values is appropriate and consistent with the strategies chosen.

3) GroupBy Aggregation (Mallikarjuna)

- Process: Categorize data into groups, calculate statistics, and check the results.

Validation: results of aggregation are correct and reproducible.

Overall results

-All three workflows demonstrated a 100% pass rate across 32 assertions.

Each workflow that was validated confirmed complete user scenarios for end-to-end integrity.

CSV roundtrip data type preservation

- Integers: 1, 2, 3, 4, 5

- Strings: ['Alice', 'Bob', 'Charlie', 'David', 'Eve'] - floats: [95.5, 87.3, 92.1, 88.7, 91.4] Booleans: [True, False, True, True, False] Dates: identified as datetime64 datatype Missing data handling strategies Methods evaluated: ffill, bfill and fillna Result: All the strategies function correctly for the scenarios which are tested.

2.7 Security Testing:

Key Takeaway: While there are potential vulnerabilities that are missed by manual review, an automated security scan can identify critical vulnerabilities that manual review may not have caught.

Using Bandit 1.9.2 Static Application Security Testing (SAST), Automation to scan the entire codebase for security vulnerabilities: 484,492 lines of Python Code across 7,218 Python Files, and this scanning took approximately 35 seconds to complete. This tool used patterns of parsing the program Abstract Syntax Tree (AST) looking for security vulnerabilities.

The result of the code analysis finished with 17,603 security vulnerabilities being found.

Table 3: Security Vulnerability Distribution by Severity

Severity	Count	Percentage
HIGH	3	0.02%
MEDIUM	72	0.41%
LOW	17,528	99.57%

Critical Vulnerabilities Identified:

Jinja2 Autoescape Disabled (HIGH) – 3 occurrences

Category: Code Injection (CWE-94)

- Location: pandas/io/formats/style.py, style_render.py

Risk: Allows execution of malicious template code

SQL Injection Vectors (MEDIUM) – 14 occurrences

Category: SQL Injection CWE-89

Location: pandas/io/sql.py

Risk: Potential compromise of the database arising from unsafe query construction

Pickle Deserialization (MEDIUM) - 29 occurrences

Category: Insecure Deserialization (CWE-502)

Location: pandas/io/pickle.py and throughout the codebase

Risk: Possible remote code execution when loading untrusted pickle files Specific remediation approaches common to all the vulnerability categories include: - Utilizing parameterized queries to minimize SQL injection attacks - Adoption of safer alternative data formats (Parquet, HDF5) as replacements for the various issues with pickle

- Jinja2 autoescape configuration recommendations

This scan identified critical security issues that could lead to code injection, SQL injection, and remote code execution. All findings are documented with specific remediation guidance.

2.8 Performance Testing

Key Takeaway: Concurrent load testing reveals breaking points not visible in single-user tests

Three performance test patterns using custom Python load generators:

1. Load Test (Sandeep): 50 concurrent users, sustained load, 1000 operations
2. Stress Test (Nithikesh): Progressive load 5→40 users until breaking point
3. Spike Test (Mallikarjuna): Traffic burst 3→30→3 users




Figure 1: Load test results showing throughput, response time distribution, CPU and memory utilization under sustained concurrent load of 50 users executing 1000 operations.

Load Test (Figure 1): - Throughput: 163.05 operations/second - Mean response: 0.0324s, p95: 0.0840s, p99: 0.1286s - Success rate: 100% (1000/1000 operations) - CPU: 39.24% average, 100% peak - Memory: 115.30 MB average, stable




Figure 2: Stress test results demonstrating catastrophic memory explosion and performance degradation at 10 concurrent users. Progressive load ramping from 5 to 40 users reveals system breaking point.

Stress Test - CRITICAL FINDING (Figure 2): - Breaking point: 10 concurrent users - Memory explosion: 5.26GB → 36.46GB (6.93x growth) - Response degradation: 10.45s → 75.61s (7.24x) - Throughput collapse: 72% drop - CPU: Saturated at 100% .




Figure 3: Spike test results showing 25x response time degradation during traffic burst (3→30→3 users) with complete recovery. Six-panel visualization includes response time, throughput, CPU

Spike Test (Figure 3): Response Time Increased from 0.0641 Seconds to 1.6214 Seconds During Spike (25.30 Times Worse) - 99th Percentile Latency Increased 19.9 Times During Spike - CPU 100% Utilized for Entire Lifecycle - Memory Increased 79.5MB During Spike and 92.7MB Retained After Spike - Recovery Time to Baseline is Perfect (0.97 Times Recovery vs Baseline). Critical Bottlenecks Are Memory Explosion, Spike Degradation and CPU Saturation - Immediate Recommendations for Memory Management Optimization, Short-Term Recommendations for Auto-Scaling and Connection Pooling, and Long-Term Recommendations for Distributed Computing (Dask).

Performance Testing Uncovered CRITICAL Issues That Nobody Saw Until They Tested Under Concurrent Load; for 10 users, this is a stopping point for Production Deployments Requiring Concurrent Access.

3. Most Significant Issues Found

3.1 Issue #1: Memory Explosion Under Stress (CRITICAL)

Author: Nithikesh Reddy (Performance Testing - Stress Test)

Test Type: Performance test with progressive load ramping.

The system shows catastrophic growth in memory, under load from concurrent users, up until a breaking point of ten concurrent users.

Technical Details (see Figure 2)

Breaking Point: 10 concurrent users

Memory Growth: 5.26 GB → 36.46 GB ≈ 6.93x increase

Response Time Degradation: from 10.45 s to 75.61 s (approximately 7.24x slower)

- Throughput Collapse: about a 72% drop in operations per second

CPU Utilization: 100% maintained throughout the stress test

As shown in Figure 2, under increasing concurrent load, the system exhibits severe resource consumption, with memory usage growing exponentially beyond the threshold of 10 users.

Implications

This behavior is the critical obstruction for any production deployment that requires concurrency, because the system cannot support even moderate levels of concurrency. It is thus unsuitable for multi-user deployments like:

- Web applications with pandas-based backend Data science platforms serving multiple users
API services that provide endpoints for data analysis - Shared computing environments

Status: Documented with immediate optimization recommendations

Remediation Recommendations:

Immediate (Critical):

Implement pooling and reuse of memory for DataFrame objects.

Implement explicit routines for garbage collection after operations.

Identify sources of leaks by profiling memory allocations.

- Set bounds on memory consumption and consider graceful degradation.

Short-term:

to implement copy-on-write semantics:

- Allow for concurrent operations by adding connection pooling.

Implement auto-scaling responsive to memory usage.

Long-term: Integrate Dask in order to provide distributed computing in the future. Provide a mechanism (e.g., file-based or streaming) for processing large datasets in a 'chunked' fashion. Revise application memory usage patterns such that they can exploit concurrent processing.

Resulting Impacts on Quality: This finding establishes that the use of concurrent multi-user environments to utilize pandas for production purposes is now substantially altered. While pandas are an excellent tool for analyzing data on single-user systems, this memory explosion behaviour would require significant architectural changes in order to make it suitable for use in a concurrent multi-user environment.

3.2 Issue #2: Jinja2 Autoescape Disabled (HIGH)

Discovered by: Sandeep Ramavath (Security Testing)

Testing Type: Static Application Security Testing (SAST) with Bandit

Description: Code injection vulnerability in HTML generation where Jinja2 template autoescape is disabled, allowing malicious template code execution.

Technical Details: - Vulnerability Type: Code Injection (CWE-94) - Severity: HIGH

- Locations: 3 instances - pandas/io/formats/style.py:3816 - pandas/io/formats/style_render.py:77 -

pandas/tests/io/formats/style/test_html.py:25 - Attack Vector: User-supplied data in DataFrame styling

Why It Matters: When pandas DataFrames are styled and rendered to HTML, user-supplied data could contain malicious template code. With autoescape disabled, this code would be executed, potentially leading to: - Cross-site

scripting (XSS) attacks - Arbitrary code execution in web applications - Data exfiltration - Session hijacking

Status: Documented with fix recommendation

Fix Recommendation:

Current (vulnerable)

```
Environment(autoescape=False)
```

Recommended fix

```
Environment(autoescape=True)
```

OR

```
Environment(autoescape=select_autoescape(['html', 'xml']))
```

Impact on System Quality: This vulnerability affects pandas' styling functionality used in Jupyter notebooks and web applications. While not critical for command-line usage, it poses significant risks in web-based data science platforms and dashboards.

3.3 Issue #3: 25x Spike Response Degradation (HIGH) :

Mallikarjuna performed a **performance testing spike test** on two applications: "Performance Testing" with simulated traffic bursts, and the second application "Web Applications with the highest level of Latency". The response time increased dramatically during the spike event (25x), although both applications recovered perfectly after the burst ended.

Technical Details (Figure 3 shows example of analysis): The Spike Pattern was 3 Users → 30 Users → 3 Users. The Response Time is as follows:

Pre-spike response time of 0.0641 seconds

During spike response time of 1.6214 seconds (25.30 times slower)

Post-spike response time of 0.0623 seconds (0.971 times baseline - perfect recovery)

p99 Latency increased 19.9 times during spike

CPU Utilization: 100% during all phases (spike pattern)

Memory Consumption: 79.5 MB during spike, 92.7 MB remained after (potential memory leak)

Throughput Maintained throughout with 0% error rate.

Figure 3 represents the temporal behaviour of the system under spike load conditions, illustrating the transient nature of performance degradation with the recoverable nature of function, and demonstrating the potential for continued memory allocation errors after.

The performance testing spike test demonstrated that the system successfully recovers from severe performance degradation. However, this performance degradation during spikes leads to poor user experience during peak usage. It impacts:

- Application of real-time data analysis
- Dashboard refresh during peak hours
- Scheduling/Staggering of Batch Jobs
- Time to complete the Auto-scaling response

Status: The spikes have been documented with a thorough analysis of CPU saturation and retention of memory.

Finding/Identified Issues: Positive finding with perfect recovery shows system resiliency; however, 25 times degradation represents severe performance degradation due to CPU saturation and allocation of memory. In addition, potential memory retention after a spike indicates a potential memory leak.

Recommendations

- Implement a system of prioritized operation queueing
- Introduce throttling of requests during peak usage periods
- Optimize the pattern of memory allocation
- Review memory residency (i.e., checking for potential leaks)
- Assess the use of async processing for operations not requiring real-time processing.

Impact on Quality of the System:

While Pandas does exhibit a strong resistance to unexpected loads, it does not exhibit graceful degradation under heavy loads. Users experience very poor performance when the system is operating at high capacity, although the reliability of the service remains very high (0% error rate).

3.4 Issue #4: SQL Injection Vectors (MEDIUM) Discovered by: Mallikarjuna (Security Testing)

Testing Type: Static Application Security Testing (SAST) with Bandit

Description: 14 instances of unsafe SQL query construction using string formatting in pandas/io/sql.py, creating SQL injection vulnerabilities.

Technical Details: - **Vulnerability Type:** SQL Injection (CWE-89) - **Severity:** MEDIUM - **Location:** pandas/io/sql.py (14 locations) - **Pattern:** String-based query construction instead of parameterized queries

Why It Matters: When pandas interfaces with SQL databases, unsafe query construction allows attackers to: - Execute arbitrary SQL commands - Bypass authentication - Read sensitive data - Modify or delete database contents - Execute database server commands

Status: Documented with parameterized query recommendations

Example Vulnerable Code:

```
# Vulnerable
query = f"SELECT * FROM {table_name} WHERE id = {user_id}"
cursor.execute(query)
```

Recommended Fix:

```
# Safe - parameterized query
query = "SELECT * FROM ? WHERE id = ?"
cursor.execute(query, (table_name, user_id))
```

Impact on System Quality: This vulnerability affects all pandas database operations. While severity is MEDIUM (requires attacker to control input), the risk is significant for web applications and data pipelines that accept user input.

4. Improvements Made to the Project

4.1 Test Coverage Improvements Before: 10% code coverage

After: 11% code coverage (+1% absolute increase)

The newly developed tests include:

Unit Tests: 15 Units for Edge Cases Not Previously Covered

Mocking/Stubbing Tests: 9 tests for isolating the behaviours of modules

Integration Tests: 6 Tests (Modules' Interactions with Each Other)

System Tests: 3 Tests: End-To-End Workflows

Performance Tests: 3 Tests (Load, Stress & Spike)

Total Number of New Tests: 48 (All 48 Tests Passed)

Test Coverage Summary:

nanops Improvements to test coverage for empty array and all NaN scenarios.

Series Constructors Added Test Coverage for Invalid Inputs and

Memory-Intensive Operations. DateTime Offsets

Enhanced Coverage for handling timezones and daylight savings times.

A 1% increase in test coverage may appear small, but it's still quite large.

Importance of the New Tests

Target important edge cases

Provide the possibility of regression protection for previously untested edge cases Document the expected behaviour of complex edge cases.

4.2 Security Vulnerability Documentation

Documented 17,603 Security Issues:

Table 1: Security Issue Remediation Status

Severity | Count | Action |

HIGH | 3 | Fix recommendations provided |

MEDIUM | 72 | Remediation guidance documented | LOW | 17,528 | Documented for awareness |

Specific Remediation Provided:

1. Jinja2 Autoescape (HIGH)

- Provided specific code fix: `Environment(autoescape=True)`
- Documented impact and attack scenarios
- Recommended testing approach

2. SQL Injection (MEDIUM)

- Provided parameterized query examples for all 14 instances
- Documented input validation strategies
- Recommended ORM usage where appropriate

3. Pickle Deserialization (MEDIUM)

- Recommended safer alternatives: Parquet, HDF5, JSON
- Documented HMAC signing approach for trusted sources
- Provided input validation patterns

Impact: Comprehensive security documentation provides pandas maintainers with: - Clear understanding of vulnerability landscape - Specific remediation code for critical issues - Prioritized action plan based on severity

4.3 Performance Bottleneck Documentation Documented 3 Critical Performance Issues:

1. Memory Explosion (CRITICAL)

- Breaking point: 10 concurrent users
- Memory growth: 6.93x (5.26GB → 36.46GB)
- Immediate action required

2. Spike Degradation (HIGH)

- Response time: 25x increase during spikes
- Perfect recovery demonstrated
- Optimization recommended

3. CPU Saturation (HIGH)

- 100% CPU usage under load
- No headroom for additional operations
- Concurrent processing redesign needed

Optimization Roadmap Created:

Immediate Actions: - Implement DataFrame memory pooling - Add explicit garbage collection - Profile memory allocations

Short-term Actions: - Implement copy-on-write semantics - Add connection pooling - Implement auto-scaling

Long-term Actions: - Consider Dask integration - Implement chunked processing - Redesign for distributed computing

Comprehensive Test Documentation Created Detailed Documentation:

1. **Test Reports** (8 comprehensive reports)
 - Unit Testing report (373 lines)
 - Mocking & Stubbing report (227 lines)
 - Mutation Testing report (285 lines)
 - Static Analysis report (103 lines)
 - Integration Testing report (125 lines)
 - System Testing report (275 lines)
 - Security Testing report (161 lines)
 - Performance Testing report (515 lines)
2. **README Files** (Execution instructions for reproducibility)
 - Unit Testing README
 - Integration Testing README
 - System Testing README
 - Security Testing README
 - Performance Testing README
3. **Summary Documents**
 - Final presentation summary (665 lines)
 - Requirements and oracles documentation

Documentation Value: - Enables reproduction of all test results - Provides context for future maintainers - Documents expected behavior for edge cases - Serves as knowledge base for testing methodologies.

5. Overall Quality Assessment

5.1 Current System Robustness

Assessment: Overall performance is Good with some Critical Performance Issues.

Strengths:

Comprehensive existing test suite, comprising thousands of tests.

Interfaces between modules are well designed and validated via integration testing.

Data integrity is ensured throughout all end-to-end workflows. No defects were found during integration and system testing. Robust error handling and input validation.

Test results:

- Integration testing: 6/6 tests passed 100%.

System testing: 3/3 workflows passed (100%).

- All newly added tests: 48/48 passed (100%).

Additional findings: • CSV round-tripping maintains all data types without loss. Missing data handling strategies work as expected. Pandas performs very well in single-user data analysis tasks, and its architecture is robust, ensuring proper functionality.

5.2 Risky Areas Identified

CRITICAL Risk - Performance: - System breaks at only 10 concurrent users with 6.93x memory explosion - Cannot handle even moderate concurrent load
- Unsuitable for multi-user production deployments without significant changes

HIGH Risk - Security: - 3 HIGH severity vulnerabilities (Jinja2 code injection)
- 72 MEDIUM severity issues (SQL injection, pickle deserialization) - Affects production deployments in web applications and shared environments

MEDIUM Risk - Performance: There is a twenty-fivefold degradation in response times during surges in traffic, with CPU reaching 100% utilization under load, leaving no headroom available for additional operations.

LOW Risk - Test Quality: Gaps identified in mutation detection undermine confidence in the overall effectiveness of testing and point to a sustained emphasis on the strength of assertions, rather than mere coverage metrics.

5.3 Maintainability and Testability Comparison

Maintainability Enhancements: - Improved exception handling by extracting messages to appropriate variables. - Better organization of imports by grouping them at the module level. - Reduced code complexity by combining nested conditionals into single-level conditionals. - Thoroughly documented all test stages for future reference.

Testability Enhancements: - Documented mocking strategies for working with external dependencies. - Utilization of deterministic testing patterns for code with time-dependent behavior. - Utilization of stronger assertions to

validate code behavior rather than just execution. - Creation of clear testing patterns for testing edge cases.

Overall, the maintainability and testability have been significantly enhanced; however, the current performance limitations will require major architectural changes in order to enable concurrent production deployments.

5.4 Most Effective Testing Methods

Performance Testing: - A critical memory explosion issue that could not be seen in the normal single-user testing was identified through performance testing. The critical point for performance testing was determined to be the number of concurrent users (the breaking point) before production-ready deployments could take place was at 10 users, with the largest impact decision to make on deployment being at this level.

Security Testing: - 17603 security vulnerabilities were found within 35 seconds, including 3 high severity vulnerabilities. Automated scanning allows for minimal effort and maximum result for security scans.

Mutation Testing: - Although good coverage was achieved through the use of coverage-based methods to assess test quality, it became apparent that there were areas of weakness in the assertion strength of the testing, and when the assertion strength was improved, the actual testing effectiveness can then be quantified.

Static Analysis: - 88 code smells were quickly identified with a high return on maintenance improvement efforts.

System Testing: - All end-to-end system workflows were able to be confirmed as functioning without any defects and confirmed the acknowledged test that quality tests have been validated and the confidence in the end user-facing features of the system has been produced.

Integration Testing: - Integration Testing could not locate any defects; it was used to confirm the quality of existing code and to validate the design of integration work.

Unit Testing (via the Coverage Expansion Method): - There was a 1% increase in code coverage with a considerable amount of time used to achieve this. Zero defects were found which is indicative of diminishing returns in analyzing mature code.

6. Lessons Learned

6.1 Sandeep Ramavath

Lesson: Test Assertions Matter More Than Coverage

My fundamental understanding of the value of test quality changed through my experiences while developing mutation testing. I first ran a set of tests against `validate_endpoints` that only detected 18% mutational coverage while at the same time “passing” every test; this led me to understand that the concept of coverage is simply a measurement of which lines of code indeed got executed and therefore does not indicate if the behaviour is validated by the tests. Upon adding more explicit assertions for return type, tuple structure, and mutual exclusivity of the left/right flags to my tests, the ability of my tests to detect mutations increased to 31%. The most important point that I gained from this experience is that the quality of a test is determined by its oracle, not the amount of code it covers.

Additionally, I learned that load testing provides an important way of establishing critical baselines. My load test showed that pandas will support up to at least 50 concurrent users (163 ops/sec) quite well; therefore, the baseline for load testing was critical in assessing what was expected out of the system under different loads. Without conducting load tests, we would have had no way of knowing whether the results from the stress tests would be a disaster or whether they were as expected.

Going forward, here is how I will apply my learning:

1. Emphasise assertion strength as I write test cases.
2. Validate test quality through mutation testing.
3. Create performance baselines through load testing before conducting stress testing.
4. Security tools can identify problems that manual inspections will not assist you.

6.2 Nithikesh Reddy (Bobbili)

Lesson: Mocking Enables Testing Impossible Scenarios

Utilizing mock filesystems means my approach to designing tests has completely changed. In the past, I thought it was always best to perform integration testing with actual files. However, mocking allows for the ability to test scenarios that would be impossible with real files. For example:

- Corrupted data that would result in a system crash
- Permission errors that require specific operating system configurations

- Missing files without creating or deleting any actual files
- All of these tests are always repeatable

The most important insight I gained through this process is that you should mock at the "seam" between your code and the outside world (i.e., the API level) when using Pandas. For example, `read_csv()` and `read_sql()`.

Another insight I gained is that stress testing is useful for finding the breaking point of your code.

My stress test showed that Pandas cannot handle more than 10 users at once before it suffers a catastrophic 633% increase in memory usage (from 5.26GB to 36.46GB). This critical issue would never be uncovered with a single-user test. Therefore, it is essential to conduct performance testing in a manner that closely approximates the true concurrent user load in order to discover performance bottlenecks related to how the application is used by real users.

Based on my experience, I will use mocking to reliably create error conditions, mock at the API boundary, conduct stress tests to determine breaking points, and monitor memory consumption during concurrent load conditions.

6.3 Mallikarjuna

Lesson: Coverage Does Not Equal Quality

My mutation detection tests for `validate_bool_kwarg` showed me that even with good coverage, I had no way to catch bugs because mutation detection showed 0%. I was surprised! What does coverage even tell us then? Coverage tells me the number of lines executed, but it doesn't tell me if my tests are validating the actual behaviour of the code. By adding testing on the boundary conditions, where type checks are done, my mutation detection now is at 10%, which showed how mutation detection really pushed me to be sure I was writing assertions that would find real bugs, not just running through paths of the code.

Additional Discovery: Information Gained from Spike Testing

My spike testing shows how pandas can take up to a 25x longer response time due to spikes in requests (0.0641 seconds to 1.6214 seconds), but will recover back to within 97% of baseline performance (0.97 vs 100%). This would not be something I could see in my normal steady-state tests.

My spike test also gave me an indication of where the possible memory leaks occurred; specifically, the spike test ran for an average of 79.5 MB of memory being used, and after the test had completed, it showed retained memory of about 92.7 MB, but when I would only test with my steady state tests, I wouldn't have seen this memory retention.

Therefore, there are several things I intend to improve upon with respect to my own understanding of what works best for testing purposes:

- 1) Use the concept of mutation testing to determine test quality.
- 2) Place more emphasis on the strength of the assertion, as opposed to the surface level on the coverage percentage.
- 3) Test the ability for an application to recover, as opposed to simply looking for peaks in performance.
- 4) Identify where you will experience memory retention due to stress testing.

7. Team Members and Roles

Sandeep Ramavath

My primary job responsibilities included performing unit tests, maintaining and expanding test coverage (i.e., adding unit tests for empty arrays, boundary conditions, and masks), creating mock and stub test cases to simulate database I/O errors (i.e., connection errors on read_sql, as well as database query errors), performing mutation tests on validate_endpoints to identify weak areas (detection increase from 18% to over 31%; over 300% improvement), performing static analysis to eliminate f-string code smell (static analysis error EM102), performing security testing to identify Jinja2 autoescape and eval vulnerability and creating a load test (50 concurrent users, 163 ops/sec baseline) to establish load testing baseline performance measures for concurrent operations, document all edge cases for nanops testing, and create deterministically created test cases for every type of failure encountered while testing database I/O.

Nithikesh Reddy (Bobbili)

Unit Tests - Unit Testing - Coverage Expansion: As part of Unit Testing for Coverage Expansion, created five tests for Series constructors that cover scenarios for the following cases: Invalid Keys, Empty Inputs, Memory Intensive Operations. Unit Testing - Unit Testing - Mocking & Stubbing: Mocking of the File System I/O Read CSV, Additionally, I have created stubs for read_excel with corrupt data and read_excel with permission error within the Mocking and Stubbing process. Mutation Testing; Created Mutation Test; Valid percentage Successful Test; Based on our tests we reached 100% detection within the Range of the Validate Percentile Function. Static Analysis; Fixed the Import Location for code smell type PLC0415. Security Testing; Identified and documented two Vulnerabilities Associated with Pickle Deserialization and insecure Temp Files. Performance Testing; Created and executed a Stress Test to determine the critical breaking point of the System. Critical breaking point of the system is the maximum users that the system was able to maintain under load; it was at 10 Users.

Major Achievements: Identified critical breaking point of System with 6.93x memory explosion. 100% mutation detection based upon the validate percentage function. Developed Tests for File System Error Conditions.

Mallikarjuna

Unit Tests; Code Conformance Testing; Unit Tests; Mocking and Stubbing Unit Tests; Unit Tests with Fuzz Test Cases; Enhanced Coverage across Timezone Edge Cases due to Migration; Unit Tests; Mutation Testing; Static Code Analysis, Including SIM102 Nested If Statements; Security Tests; Performance Tests.

Accomplished 25x Performance Degradation Spike and Complete Recovery; Documented 14 SQL Injection Vulnerabilities and Provided Mitigation Strategies; Established Deterministic Time-Series Calculations, Including Unit Testing of Time-Series (with Date Time and XML Parsing), Using Mocked and Stubbed Operations.